

## 1. Definition:

Elasticsearch is an open source, full-text search and analysis engine, based on the Apache Lucene search engine.

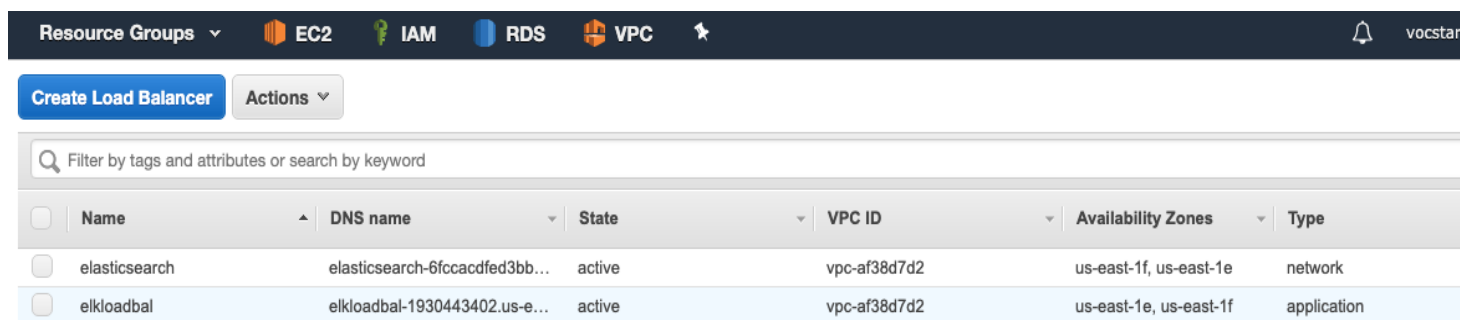
Logstash is a log aggregator that collects data from various input sources, executes different transformations and enhancements and then ships the data to various supported output destinations.

Kibana is a visualization layer that works on top of Elasticsearch, providing users with the ability to analyze and visualize the data.

Beats are lightweight agents that are installed on edge hosts to collect different types of data for forwarding into the stack.

## 2. AWS Service Setup

At the heart of our AWS setup are two load balancers








<input type="checkbox"/>	Name	DNS name	State	VPC ID	Availability Zones	Type
<input type="checkbox"/>	elasticsearch	elasticsearch-6fccacdfed3bb...	active	vpc-af38d7d2	us-east-1f, us-east-1e	network
<input type="checkbox"/>	elkloadbal	elkloadbal-1930443402.us-e...	active	vpc-af38d7d2	us-east-1e, us-east-1f	application

There is network (elasticsearch) and application (kibana) that were used. Both achieved the same end result; thus, it was a matter of exploring the options that were available.

Below is a screenshot of the elasticsearch load balancer settings:

## Jerome Thompson – ES Stack

<b>Name</b>	elasticsearch
<b>ARN</b>	arn:aws:elasticloadbalancing:us-east-1:765128034604:loadbalancer/net/elasticsearch/6fccacdfed3bb83b 
<b>DNS name</b>	elasticsearch-6fccacdfed3bb83b.elb.us-east-1.amazonaws.com  (A Record)
<b>State</b>	active
<b>Type</b>	network
<b>Scheme</b>	internet-facing
<b>IP address type</b>	ipv4
<b>VPC</b>	<a href="#">vpc-af38d7d2</a> 
<b>Availability Zones</b>	<a href="#">subnet-842ebb8a - us-east-1f</a>  IPv4 address: Assigned by AWS  <a href="#">subnet-08c02639 - us-east-1e</a>  IPv4 address: Assigned by AWS  <a href="#">Edit subnets</a>
<b>Hosted zone</b>	Z26RNL4JYFTOTI
<b>Creation time</b>	July 18, 2020 at 4:07:06 PM UTC-4

### Attributes






<b>Deletion protection</b>	Disabled
<b>Cross-Zone Load Balancing</b>	Disabled
<b>Access logs</b>	Disabled

The following snippet is of the elasticsearch listener and target group

<input type="checkbox"/>	Listener ID	Security policy	SSL Certificate	ALPN policies	Default action
<input type="checkbox"/>	<b>TCP : 9200</b> arn...cc3e6d5fa8578f4f ▾	N/A	N/A	N/A	Forward to <a href="#">estarget</a>
<input checked="" type="checkbox"/>	<b>TCP : 9300</b> arn...a13393714e4710f9 ▾	N/A	N/A	N/A	Forward to <a href="#">elasticsearch93</a>

Below is a screenshot of the kibana load balancer settings:

## Jerome Thompson – ES Stack

<b>Name</b>	elkloadbal
<b>ARN</b>	arn:aws:elasticloadbalancing:us-east-1:765128034604:loadbalancer/app/elkloadbal/56d01e7fc64091e6 
<b>DNS name</b>	elkloadbal-1930443402.us-east-1.elb.amazonaws.com  (A Record)
<b>State</b>	active
<b>Type</b>	application
<b>Scheme</b>	internet-facing
<b>IP address type</b>	ipv4 <a href="#">Edit IP address type</a>
<b>VPC</b>	<a href="#">vpc-af38d7d2</a> 
<b>Availability Zones</b>	<a href="#">subnet-08c02639 - us-east-1e</a>  IPv4 address: Assigned by AWS  <a href="#">subnet-842ebb8a - us-east-1f</a>  IPv4 address: Assigned by AWS  <a href="#">Edit subnets</a>
<b>Hosted zone</b>	Z35SXDOTRQ7X7K
<b>Creation time</b>	July 15, 2020 at 3:05:55 AM UTC-4





### Security

**Security groups** [sg-03d261985cd0b395d](#), [es-load-balancer-wizard-1](#)  
• load-balancer-wizard-1 created on 2020-07-15T02:58:03.314-04:00

Below is a shot of the kibana listeners / target group

<input type="checkbox"/>	Listener ID	Security policy	SSL Certificate	Rules
<input type="checkbox"/>	<b>HTTP : 5601</b> arn...2e034d0f4baddb64 ▾	N/A	N/A	Default: forwarding to <a href="#">kibanatrggrp</a> <a href="#">View/edit rules</a>

As follows we note the target groups (there are four but we are only using 2 of them)

<input type="checkbox"/>	Name	ARN	Port	Protocol	Target type	Load balancer	VPC ID
<input type="checkbox"/>	<a href="#">elasticsearch93</a>	 arn:aws:elasticload...	9300	TCP	Instance	elasticsearch	vpc-af38d7d2
<input type="checkbox"/>	<a href="#">estarget</a>	 arn:aws:elasticload...	9200	TCP	Instance	elasticsearch	vpc-af38d7d2
<input type="checkbox"/>	<a href="#">estrgrp</a>	 arn:aws:elasticload...	9200	HTTP	Instance		vpc-af38d7d2
<input type="checkbox"/>	<a href="#">kibanatrggrp</a>	 arn:aws:elasticload...	5601	HTTP	Instance	elkloadbal	vpc-af38d7d2

We are making use of the elasticsearch group that monitors traffic for port 9200

estarget

Delete

arn:aws:elasticloadbalancing:us-east-1:765128034604:targetgroup/estarget/63e9cf10ee2bc054

Basic configuration

Target type Instance	Protocol : Port TCP : 9200	VPC <a href="#">vpc-af38d7d2</a>	Load balancer <a href="#">elasticsearch</a>
-------------------------	-------------------------------	-------------------------------------	------------------------------------------------

Group details

Targets

Monitoring

Tags

Health check settings

Edit

Protocol HTTP	Unhealthy threshold 3
Path /	Timeout 6
Port traffic-port	Interval 30
Healthy threshold 3	Success codes 200-399

Along with the instance tied to that node (note that we are using one node). Based on “The simple answer is that **Elasticsearch** is designed specifically to not **need a load balancer.**” ([How to Connect to Multiple Elasticsearch Client Nodes Without a Load Balancer](#), November 2017). So only went with one node and keep the option of the load balancer so that I could use the name when doing elasticsearch setups.

<input type="checkbox"/>	Instance ID	Name	Port	Zone	Status	Status details
<input type="checkbox"/>	<a href="#">i-0f62f5f221599de82</a>	elasticsearch	9200	us-east-1f	<span>healthy</span>	

Next we see a picture of the Kibana Load balancer kibanatrggrp

kibanatrggrp

Delete

arn:aws:elasticloadbalancing:us-east-1:765128034604:targetgroup/kibanatrggrp/534d318f9b4944bd

Basic configuration

Target type Instance	Protocol : Port HTTP : 5601	VPC <a href="#">vpc-af38d7d2</a>	Load balancer <a href="#">elkloadbal</a>
-------------------------	--------------------------------	-------------------------------------	---------------------------------------------

Group details

Targets

Monitoring

Tags

Health check settings

Edit

Protocol HTTP	Unhealthy threshold 2
Path /	Timeout 5
Port traffic-port	Interval 30
Healthy threshold 5	Success codes 302










Also note the setup for its instance’s association

<input type="checkbox"/>	Instance ID	Name	Port	Zone	Status	Status details
<input type="checkbox"/>	<a href="#">i-08b3c91038e0ba78a</a>	kibana	5601	us-east-1f	<span>healthy</span>	
<input type="checkbox"/>	<a href="#">i-06d839d89b5d794fb</a>	kibana	5601	us-east-1e	<span>healthy</span>	

Next we take a look at the launch configurations, which uses AMIs to define what OS and specs should be spun up when auto scaling

Create launch configuration						Create Auto Scaling group	Copy to launch template	Actions
Filter: <input type="text" value="Filter launch configurations..."/>								
<input type="checkbox"/>	Name	AMI ID	Instance Type	Spot Price	Creation Time			
<input type="checkbox"/>	kibana_launch...	ami-0e4ebb21fbe0dcc05	t2.micro		July 18, 2020 at 8:14:42 PM UT...			
<input type="checkbox"/>	kib_launch_conf	ami-0e4ebb21fbe0dcc05	t2.micro		July 18, 2020 at 7:47:42 PM UT...			
<input type="checkbox"/>	es_launch_conf1	ami-078ffe7ed441d34c5	t2.small		July 18, 2020 at 5:11:32 PM UT...			

Below are some of the AMIs I created to make things must simpler when do the setups

Quick Start							1 to 9 of 9 AMIs	
My AMIs		linux_apachev2 - ami-0031d87e99496d71b					Select	64-bit (x86)
AWS Marketplace								
Community AMIs		winsisvol - ami-0075aea1002d5cbfc					Select	64-bit (x86)
Ownership								
<input checked="" type="checkbox"/> Owned by me								
<input type="checkbox"/> Shared with me								
Architecture								
<input type="checkbox"/> 32-bit (x86)								
<input type="checkbox"/> 64-bit (x86)								
<input type="checkbox"/> 64-bit (Arm)								
Root device type								
<input type="checkbox"/> EBS								
<input type="checkbox"/> Instance store								
		es_kib_ami - ami-039768cdd8e394650					Select	64-bit (x86)
		linux_apache - ami-04dfb361fb947c60					Select	64-bit (x86)
		elasticsearch - ami-078ffe7ed441d34c5					Select	64-bit (x86)
		kibanaUsed - ami-0832231730aee8cd1					Select	64-bit (x86)
		windowsis - ami-0886ecd5bdcf903d3					Select	64-bit (x86)
		kibana - ami-0de1de68e59fa4cc0					Select	64-bit (x86)
		kibanav2 - ami-0e4ebb21fbe0dcc05					Select	64-bit (x86)

I create a lot of AMIs for different purposes

The last key service for the auto scaling group, this what governs how the instances are scaled per target group. Notice that elasticsearch remains at one will the kibana starts at 2 and can grow to 4.

Create Auto Scaling groupActions

Filter:

<input type="checkbox"/>	Name	Launch Configuration	Instances	Desired	Min	Max	Availability Zones	Default Cooldown	Health Check Gr
<input type="checkbox"/>	KIB Auto Scali...	kibana_launch_conf	2	2	2	4	us-east-1e, us-east-1f	300	300
<input type="checkbox"/>	ES Auto Scalin...	es_launch_conf1	1	1	1	1	us-east-1f	300	300

The scaling policy is note below for kibana, when the CPU > 85 scaling occurs

## Scale Group Size

**Policy type:** Target Tracking scaling

**Execute policy when:** As required to maintain Average CPU Utilization at 85

**Take the action:** Add or remove instances as required

**Instances need:** 300 seconds to warm up after scaling

**Disable scale-in:** No

Below we take of instances that are in play

	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)	IPv4 Public IP	IPv6 IPs
		i-0b2febb397425c8a8	t2.small	us-east-1e	stopped		None		-	-
	linux_hhtpd	i-03e56a22f288ebfda	t2.micro	us-east-1a	running	2/2 checks passed	None	ec2-34-207-201-161.co...	34.207.201.161	-
	winsilis	i-02d45d2ec2fc19ee5	t2.micro	us-east-1e	running	2/2 checks passed	None	ec2-34-229-161-82.co...	34.229.161.82	-
	kibana	i-08b3c91038e0ba78a	t2.micro	us-east-1f	running	2/2 checks passed	None	ec2-34-239-156-10.co...	34.239.156.10	-
	elasticsearch	i-0f62f5f221599de82	t2.small	us-east-1f	running	2/2 checks passed	None	ec2-3-236-52-168.com...	3.236.52.168	-
	kibana	i-06d839d89b5d794fb	t2.micro	us-east-1e	running	2/2 checks passed	None	ec2-18-208-119-107.co...	18.208.119.107	-

### 3. Server Setup

#### 1. Linux Apache setup

Below are the aws settings

linux_hhtpd	i-03e56a22f288ebfda	t2.micro	us-east-1a	running	2/2 checks passed	None	ec2-34-207-201-161.co...	34.207.201.161
Instance ID	i-03e56a22f288ebfda					Public DNS (IPv4)	ec2-34-207-201-161.compute-1.amazonaws.com	
Instance state	running					IPv4 Public IP	34.207.201.161	
Instance type	t2.micro					IPv6 IPs	-	
Finding	You may not have permission to access AWS Compute Optimizer.					Elastic IPs	-	
Private DNS	ip-172-31-34-218.ec2.internal					Availability zone	us-east-1a	
Private IPs	172.31.34.218					Security groups	<a href="#">es_kib_sg</a> , <a href="#">view inbound rules</a> , <a href="#">view outbound rules</a>	
Secondary private IPs	-					Scheduled events	No scheduled events	
VPC ID	vpc-af38d7d2					AMI ID	<a href="#">linux_apache (ami-04dfb361fb9f47c60)</a>	
Subnet ID	subnet-1497484b					Platform details	Linux/UNIX	
Network interfaces	<a href="#">eth0</a>					Usage operation	RunInstances	
IAM role	-					Source/dest. check	True	
Key pair name	elk_key					T2/T3 Unlimited	Disabled	
Owner	765128034604					EBS-optimized	False	
Launch time	July 18, 2020 at 8:51:25 PM UTC-4 (54 hours)					Root device type	ebs	
Termination protection	False					Root device	<a href="#">/dev/xvda</a>	
Lifecycle	normal					Block devices	<a href="#">/dev/xvda</a>	
Monitoring	basic					Elastic Graphics ID	-	
Alarm status	None					Elastic Inference accelerator ID	-	
Kernel ID	-					Capacity Reservation	-	
RAM disk ID	-					Capacity Reservation Settings	Open	
Placement group	-					Outpost Arm	-	
Partition number	-							
Virtualization	hvm							
Reservation	r-02a940a6be166f298							
AMI launch index	0							
Tagging	default							

Below was used to yum install filebeat

```
[root@ip-172-31-34-218 ~]# cat /etc/yum.repos.d/elastic-beats.repo
[elastic-6.x]
name=Elastic repository for 6.x packages
baseurl=https://artifacts.elastic.co/packages/oss-6.x/yum
gpgcheck=1
gpgkey=https://artifacts.elastic.co/GPG-KEY-elasticsearch
enabled=1
autorefresh=1
type=rpm-md
[root@ip-172-31-34-218 ~]#
```

```
- type: log

# Change to true to enable this input configuration.
enabled: true

# Paths that should be crawled and fetched. Glob based paths.
paths:
  - /var/log/messages
  - /var/log/httpd/*
  - /var/log/filebeat/*
  #- c:\programdata\elasticsearch\logs\*
```

Above shows the path that will be used filebeat to feed elasticsearch

```
===== Dashboards =====
# These settings control loading the sample dashboards to the Kibana index. Loading
# the dashboards is disabled by default and can be enabled either by setting the
# options here, or by using the `--setup` CLI flag or the `setup` command.
setup.dashboards.enabled: true

# The URL from where to download the dashboard archive. By default this URL
# has a value which is computed based on the Beat name and version. For released
# versions, this URL points to the dashboard archive on the artifacts.elastic.co
# website.
#setup.dashboards.url:
```

Above shows that dashboards have been sent to kibana server



## Jerome Thompson – ES Stack

```
#===== Kibana =====  
# Starting with Beats version 6.0.0, the dashboards are loaded via the Kibana API.  
# This requires a Kibana endpoint configuration.  
setup.kibana:  
  
# Kibana Host  
# Scheme and port can be left out and will be set to the default (http and 5601)  
# In case you specify an additional path, the scheme is required: http://localhost:5601/pa  
# IPv6 addresses should always be defined as: https://[2001:db8::1]:5601  
host: "elkloadbal-1930443402.us-east-1.elb.amazonaws.com:5601"  
  
# Kibana Space ID  
# ID of the Kibana Space into which the dashboards should be loaded. By default,  
# the Default Space will be used.  
#space.id:
```

Above shows the kibana settings

```
#----- Elasticsearch output -----  
output.elasticsearch:  
# Array of hosts to connect to.  
hosts: ["elasticsearch-6fccacdfed3bb83b.elb.us-east-1.amazonaws.com:9200"]  
  
# Enabled ilm (beta) to use index lifecycle management instead daily indices.  
#ilm.enabled: false  
  
# Optional protocol and basic auth credentials.  
#protocol: "https"  
#username: "elastic"  
#password: "changeme"
```

Above shows the elasticsearch settings

```
#===== Logging =====  
# Sets log level. The default log level is info.  
# Available log levels are: error, warning, info, debug  
logging.level: debug  
logging.to_files: true  
logging.files:  
  path: /var/log/filebeat  
  name: filebeat  
  keepfiles: 7  
  permissions: 0644
```

Above shows the logging option chosen

```
vim /etc/filebeat/modules.d/  
filebeat modules enable apache2  
systemctl restart filebeat
```

Above shows the apache2 being enabled

Below shows the filebeat options for starting and enabling:

```
systemctl start filebeat  
systemctl enable filebeat
```





## 2. Windows 98 IIS Server setup

### 1. Filebeat

Can be used to install filebeat

<https://www.elastic.co/guide/en/beats/filebeat/current/filebeat-installation.html>

filebeat.inputs:

# Each - is an input. Most options can be set at the input level, so  
# you can use different inputs for various configurations.  
# Below are the input specific configurations.

- type: log

# Change to true to enable this input configuration.  
enabled: false

# Paths that should be crawled and fetched. Glob based paths.  
paths:

#- C:\inetpub\logs\LogFiles\W3SVC1  
#- /var/log/\*.log  
#- c:\programdata\elasticsearch\logs\\*

Above shows the input for filebeat

```
# ===== Dashboards =====
# These settings control loading the sample dashboards to the Kibana index. Loading
# the dashboards is disabled by default and can be enabled either by setting the
# options here or by using the `setup` command.
setup.dashboards.enabled: true

# The URL from where to download the dashboards archive. By default this URL
# has a value which is computed based on the Beat name and version. For released
# versions, this URL points to the dashboard archive on the artifacts.elastic.co
# website.
#setup.dashboards.url:
```

Above shows the option for sending dashboards to kibana has been enabled

```
# ===== Kibana =====

# Starting with Beats version 6.0.0, the dashboards are loaded via the Kibana API.
# This requires a Kibana endpoint configuration.
setup.kibana:

# Kibana Host
# Scheme and port can be left out and will be set to the default (http and 5601)
# In case you specify an additional path, the scheme is required: http://localhost:5601/path
# IPv6 addresses should always be defined as: https://[2001:db8::1]:5601
host: "elkloadbal-1930443402.us-east-1.elb.amazonaws.com:5601"

# Kibana Space ID
# ID of the Kibana Space into which the dashboards should be loaded. By default,
# the Default Space will be used.
#space.id:
```

Above shows the kibana host setup settings

```
# ----- Elasticsearch Output -----
output.elasticsearch:
# Array of hosts to connect to.
hosts: ["elasticsearch-6fccacdfed3bb83b.elb.us-east-1.amazonaws.com:9200"]

# Protocol - either `http` (default) or `https`.
#protocol: "https"

# Authentication credentials - either API key or username/password.
#api_key: "id:api_key"
#username: "elastic"
#password: "changeme"
```

Above shows the elasticsearch settings setup

```
# ===== Logging =====
```

```
# Sets log level. The default log level is info.
# Available log levels are: error, warning, info, debug
#logging.level: debug
logging.level: debug
logging.to_files: true
logging.files:
  path: C:\Program Files\filebeat-7\Logs
  name: filebeat
  keepfiles: 7
  permissions: 0644
# At debug level, you can selectively enable logging only for some components.
Above shows the logging settings
```



[Above shows the what is displayed in kibana for filebeat](#)

## 2. Winlogbeat

Can be used to install winlogbeat

<https://www.elastic.co/guide/en/beats/winlogbeat/current/winlogbeat-installation.html>

```
#===== Winlogbeat specific options =====
```

```
# event_logs specifies a list of event logs to monitor as well as any  
# accompanying options. The YAML data type of event_logs is a list of  
# dictionaries.
```

```
#
```

```
# The supported keys are name (required), tags, fields, fields_under_root,  
# forwarded, ignore_older, level, event_id, provider, and include_xml. Please  
# visit the documentation for the complete details of each option.
```

```
# https://go.es.io/WinlogbeatConfig
```

```
winlogbeat.event_logs:
```

- name: Application  
 ignore\_older: 72h
- name: Security
- name: System
- name: ForwardedEvents  
 tags: [forwarded]

```
#===== Elasticsearch template setting =====
```

Above shows the events that will be sent to kibana

```
#===== Dashboards =====
```

```
# These settings control loading the sample dashboards to the Kibana:
```

```
#index. Loading
```

```
# the dashboards is disabled by default and can be enabled either by setting the  
# options here, or by using the `setup` CLI flag or the `setup` command.
```

```
setup.dashboards.enabled: true
```

```
# The URL from where to download the dashboards archive. By default this URL  
# has a value which is computed based on the Beat name and version. For released  
# versions, this URL points to the dashboard archive on the artifacts.elastic.co  
# website.
```

```
#setup.dashboards.url:
```

Above shows the dashboard is being sent to kibana

```
#===== Outputs =====
```

```
# Configure what output to use when sending the data collected by the beat.
```

```
#----- Elasticsearch output -----
```

```
output.elasticsearch:
```

```
# Array of hosts to connect to.
```

```
hosts: ["elasticsearch-6fccacdfed3bb83b.elb.us-east-1.amazonaws.com:9200"]
```

```
# Enabled ilm (beta) to use index lifecycle management instead daily indices.
```

```
#ilm.enabled: false
```

```
# Optional protocol and basic auth credentials.
```

```
#protocol: "https"
```

```
#username: "elastic"
```

```
#password: "changeme"
```

Above shows the elasticsearch settings for winlogbeat

```
#===== Logging =====

# Sets log level. The default log level is info.
# Available log levels are: error, warning, info, debug
#logging.level: debug
logging.level: debug
logging.to_files: true
logging.files:
  path: C:\Program Files\winlogbeat-7\Logs
  name: winlogbeat
  keepfiles: 7
  permissions: 0644

# At debug level, you can selectively enable logging only for some components
Above shows the logging for winlogbeat
```



Above shows the services have been started



Above shows the kibana display for winlogbeat

### 3. Elasticsearch server setup

Used the same repo noted in setup of filebeat on Linux Apache server to yum install elasticsearch.

I followed the posted guide for elasticsearch install (tried a couple things that didn't work out such as index renaming).



```
# ----- Paths -----  
#  
# Path to directory where to store the data (separate multiple locations by comma):  
#  
path.data: /var/lib/elasticsearch  
#  
# Path to log files:  
#  
path.logs: /var/log/elasticsearch  
#  
# ----- Memory -----
```

Above shows the paths for elasticsearch settings

```
#  
# ----- Network -----  
#  
# Set the bind address to a specific IP (IPv4 or IPv6):  
#  
#network.host: 192.168.0.1  
network.host: 0.0.0.0  
network.bind_host: 0.0.0.0  
network.publish_host: 0.0.0.0  
transport.host: localhost  
transport.tcp.port: 9300  
#  
# Set a custom port for HTTP:  
#  
http.port: 9200  
#  
# For more information, consult the network module documentation.  
#  
# ----- Discovery -----
```

Above shows the network config for elasticsearch

```
# ----- Discovery -----  
#  
# Pass an initial list of hosts to perform discovery when this node is started:  
# The default list of hosts is ["127.0.0.1", "[:,1]"]  
#  
#discovery.seed_hosts: ["host1", "host2"]  
#  
# Bootstrap the cluster using an initial set of master-eligible nodes:  
#  
#cluster.initial_master_nodes: ["node-1", "node-2"]  
#  
# For more information, consult the discovery and cluster formation module documentation.  
#  
cluster.name: elasticsearch  
discovery.type: single-node  
discovery.seed_providers: ec2  
discovery.ec2.endpoint: ec2.us-east-1.amazonaws.com  
discovery.ec2.tag.Name: "elasticsearch"  
http.cors.enabled: true  
http.cors.allow-origin: "*"
```

Above shows the discovery option for elasticsearch as it regards the ec2 instances



```
[root@ip-172-31-69-138 ~]# tail -f /var/log/elasticsearch/elasticsearch.log
[2020-07-21T04:07:30.247][INFO ][o.e.c.m.MetadataMappingService] [ip-172-31-69-138.ec2.internal] [winlogbeat-7.8.0-2020.07.21/s3whThUHTCSQS-CpXo_A1]
A] update_mapping [_doc]
[2020-07-21T04:07:39.309][INFO ][o.e.c.m.MetadataMappingService] [ip-172-31-69-138.ec2.internal] [winlogbeat-7.8.0-2020.07.21/s3whThUHTCSQS-CpXo_A1]
A] update_mapping [_doc]
[2020-07-21T04:08:02.416][INFO ][o.e.c.m.MetadataMappingService] [ip-172-31-69-138.ec2.internal] [winlogbeat-7.8.0-2020.07.21/s3whThUHTCSQS-CpXo_A1]
A] update_mapping [_doc]
```

Above shows the elasticsearch logs after being started

Below can be used to start and enable elasticsearch:

systemctl enable elasticsearch

systemctl start elasticsearch

Below depicts how to delete an index for the command line:

/usr/share/elasticsearch/bin/elasticsearch-plugin install discovery-ec2

curl -XDELETE 'http://http://ec2-54-160-96-119.compute-1.amazonaws.com/:9200/.kibana\_1' -  
-header "content-type: application/JSON"

#### 4. Kibana server setup

Used the same repo noted in setup of filebeat on Linux apache server to yum install kibana-oss

```
# Kibana is served by a back end server. This setting specifies the port to use.
server.port: 5601

# Specifies the address to which the Kibana server will bind. IP addresses and host names
# The default is 'localhost', which usually means remote machines will not be able to con
# To allow connections from remote users, set this parameter to a non-loopback address.
#server.host: "localhost"
server.host: "0.0.0.0"

# Enables you to specify a path to mount Kibana at if you are running behind a proxy.
# Use the `server.rewriteBasePath` setting to tell Kibana if it should remove the basePat
# from requests it receives, and to prevent a deprecation warning at startup.
# This setting cannot end in a slash.
#server.basePath: ""
```

Above shows Part 1 of the kibana server settings local port and server ip settings

```
# The Kibana server's name. This is used for display purposes.
#server.name: "your-hostname"
server.name: "kibana"

# The URLs of the Elasticsearch instances to use for all your queries.
elasticsearch.hosts: ["http://elasticsearch-6fccacdfed3bb83b.elb.us-east-1.amazonaws.com:9200"]

# When this setting's value is true Kibana uses the hostname specified in the server.host
# setting. When the value of this setting is false, Kibana uses the hostname of the host
# that connects to this Kibana instance.
#elasticsearch.preserveHost: true
elasticsearch.preserveHost: false

# Kibana uses an index in Elasticsearch to store saved searches, visualizations and
# dashboards. Kibana creates a new index if the index doesn't already exist.
kibana.index: ".kibana"

# The default application to load.
#kibana.defaultAppId: "home"
```

Above shows Part2 of kibana settings with additionally server info

```
#elasticsearch.requestTimeout: 30000
elasticsearch.requestTimeout: 60000

# List of Kibana client-side headers to send to Elasticsearch. To send
# headers, set this value to [] (an empty list).
#elasticsearch.requestHeadersWhitelist: [ authorization ]

# Header names and values that are sent to Elasticsearch. Any custom he
# by client-side headers, regardless of the elasticsearch.requestHeader
#elasticsearch.customHeaders: {}

# Time in milliseconds for Elasticsearch to wait for responses from sha
#elasticsearch.shardTimeout: 30000

# Time in milliseconds to wait for Elasticsearch at Kibana startup befor
#elasticsearch.startupTimeout: 5000

# Logs queries sent to Elasticsearch. Requires logging.verbose set to t
#elasticsearch.logQueries: false
elasticsearch.logQueries: true

# Specifies the path where Kibana creates the process ID file.
#pid.file: /var/run/kibana.pid

# Enables you specify a file where Kibana stores log output.
#logging.dest: stdout
logging.dest: /var/log/kibana/kibana.log
```

Above shows the Part 3 of elasticsearch settings

Below can be used to stop and start server:

systemctl enable kibana

systemctl start kibana

#### 4. Reports and Dashboards

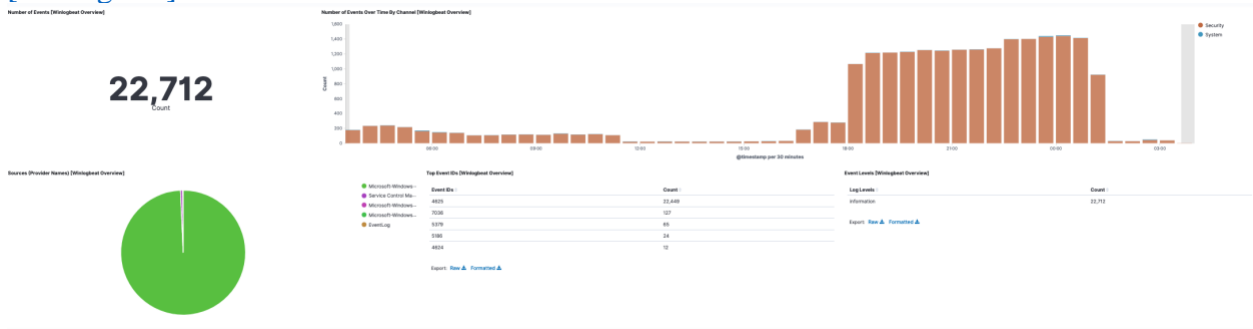
##### 1. Dashboards

[Access and error logs ECS](#)

## Jerome Thompson – ES Stack

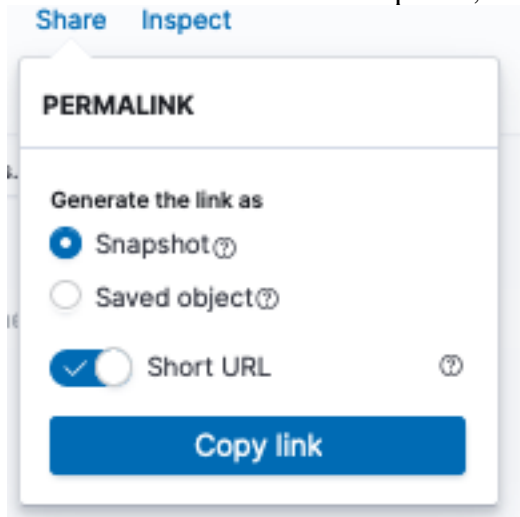


## [Winlogbeat] Overview



## 2. Reports

I do not have the save to csv options, so I have hyper-linked reports below



[windows\\_iis\\_filebeat7\\_iis\\_error](#)

is used to display windows iis error related entries

[windows\\_winlogbeat\\_admin\\_activity](#)

is used to display admin related windows events

## 5. Conclusion

This was an amazing project that showed how to harness the power of the ES Stack through creation of servers that did various pieces of the stack along with ones that feed logs to elasticsearch.

## 6. Reference

How to Connect to Multiple Elasticsearch Client Nodes Without a Load Balancer, ObjectRocket, November 11 2017, <https://www.objectrocket.com/blog/elasticsearch/elasticsearch-and-load-balancers>,